

[!\[\]\(919a2cb85b99741a73c0c31a427236a8\_img.jpg\) Home](#) [!\[\]\(c9cd5a1c35167a83f09a35036fe5dcbd\_img.jpg\) Cookbook](#) [Vector stores](#) [Faiss](#)

# LangChain with FAISS Vector DB

Example by Joselin James. Example was adapted to use README.md as the source of documents in the DB.

```
import os  
  
os.environ["TRULENS_OTEL_TRACING"] = "0"
```



## Import packages

```
# !pip install trulens trulens-apps-langchain faiss-cpu unstructured
```



```
from typing import List  
  
from langchain.callbacks.manager import CallbackManagerForRetrieverRun  
from langchain.chains import ConversationalRetrievalChain  
from langchain.chat_models import ChatOpenAI  
from langchain.document_loaders import UnstructuredMarkdownLoader  
from langchain.embeddings.openai import OpenAIEmbeddings  
from langchain_core.documents import Document  
from langchain.text_splitter import CharacterTextSplitter  
from langchain.vectorstores import FAISS  
from langchain.vectorstores.base import VectorStoreRetriever  
import nltk  
import numpy as np  
from trulens.core import Feedback  
from trulens.core import Select  
from trulens.core import TruSession  
from trulens.apps.langchain import TruChain
```



## Set API keys

```
import os

if "OPENAI_API_KEY" not in os.environ:
    os.environ["OPENAI_API_KEY"] = "sk-..."
```



## Create vector db

```
# Create a local FAISS Vector DB based on README.md .
loader = UnstructuredMarkdownLoader("README.md")
nltk.download("averaged_perceptron_tagger")
documents = loader.load()

text_splitter = CharacterTextSplitter(chunk_size=1000, chunk_overlap=0)
docs = text_splitter.split_documents(documents)

embeddings = OpenAIEMBEDDINGS()
db = FAISS.from_documents(docs, embeddings)

# Save it.
db.save_local("faiss_index")
```



## Create retriever

```
class VectorStoreRetrieverWithScore(VectorStoreRetriever):
    def _get_relevant_documents(
        self, query: str, *, run_manager: CallbackManagerForRetrieverRun
    ) -> List[Document]:
        if self.search_type == "similarity":
            docs_and_scores = (
                self.vectorstore.similarity_search_with_relevance_scores(
                    query, **self.search_kwargs
                )
            )

        print("From relevant doc in vec store")
        docs = []
        for doc, score in docs_and_scores:
            if score > 0.6:
                doc.metadata["score"] = score
                docs.append(doc)
        elif self.search_type == "mmr":
            docs = self.vectorstore.max_marginal_relevance_search(
                query, **self.search_kwargs
            )
        else:
            raise ValueError(f"search_type of {self.search_type} not allowed.")
    return docs
```



## Create app

```
# Create the example app.
class FAISSWithScore(FAISS):
    def as_retriever(self) -> VectorStoreRetrieverWithScore:
        return VectorStoreRetrieverWithScore(
            vectorstore=self,
            search_type="similarity",
            search_kwargs={"k": 4},
        )

class FAISSStore:
    @staticmethod
    def load_vector_store():
        embeddings = OpenAIEMBEDDINGS()
        faiss_store = FAISSWithScore.load_local(
            "faiss_index", embeddings, allow_dangerous_deserialization=True
        )
        print("Faiss vector DB loaded")
        return faiss_store
```

## Set up evals

```
from trulens.providers.openai import OpenAI

# Create a feedback function.
openai = OpenAI()

f_context_relevance = (
    Feedback(openai.context_relevance, name="Context Relevance")
    .on_input()
    .on(
        Select.Record.app.combine_docs_chain._call.args.inputs.input_documents[
            :
        ].page_content
    )
    .aggregate(np.min)
)
```

```
# Bring it all together.

def load_conversational_chain(vector_store):
    llm = ChatOpenAI(
        temperature=0,
        model_name="gpt-4",
    )
    retriever = vector_store.as_retriever()
    chain = ConversationalRetrievalChain.from_llm(
        llm, retriever, return_source_documents=True
    )

    # workaround to avoid hitting ValueError: run not supported when there is not
    # exactly one output key. Got ['answer', 'source_documents'] in langchain/chains/
    # base.py:546, in Chain._run_output_key(self)
    chain.return_source_documents = False
    truchain = TruChain(chain, feedbacks=[f_context_relevance], with_hugs=False)
    chain.return_source_documents = True
    return chain, truchain
```

```
# Run example:
vector_store = FAISSStore.load_vector_store()
chain, tru_chain_recorder = load_conversational_chain(vector_store)

with tru_chain_recorder as recording:
    ret = chain({"question": "What is trulens?", "chat_history": ""})
```

```
# Check result.
ret
```

```
# Check that components of the app have been instrumented despite various
# subclasses used.
tru_chain_recorder.print_instrumented()
```

```
# Start dashboard to inspect records.
TruSession().run_dashboard()
```